

## 01 SQL Interview Questions & Answers

### SQL Interview Questions & Answers

#### Basic SQL Questions

#### DDL & DML

**What is the difference between DDL and DML?  
Provide examples of each.**

**DDL (Data Definition Language)** is used to define and modify database structure:

- **Purpose:** Create, modify, or delete database objects
- **Examples:** CREATE, ALTER, DROP, TRUNCATE
- **Auto-commit:** DDL statements are automatically committed

**DML (Data Manipulation Language)** is used to manipulate data within tables:

- **Purpose:** Insert, update, delete, or retrieve data
- **Examples:** INSERT, UPDATE, DELETE, SELECT
- **Transaction control:** Can be rolled back using ROLLBACK

-- DDL Examples

```
CREATE TABLE employees (  
    emp_id INT PRIMARY KEY,  
    name VARCHAR(50),  
    salary DECIMAL(10,2)  
);
```

```
ALTER TABLE employees ADD COLUMN department VARCHAR(30);  
DROP TABLE temp_table;
```

-- DML Examples

```
INSERT INTO employees VALUES (1, 'John Doe', 50000);  
UPDATE employees SET salary = 55000 WHERE emp_id = 1;  
DELETE FROM employees WHERE emp_id = 1;  
SELECT * FROM employees;
```

#### How do you rename a table in SQL?

-- Standard SQL (MySQL, PostgreSQL)

```
ALTER TABLE old_table_name RENAME TO new_table_name;
```

-- SQL Server

```
EXEC sp_rename 'old_table_name', 'new_table_name';
```

```
-- Oracle
ALTER TABLE old_table_name RENAME TO new_table_name;
```

**Write a query to rename a column in a table.**

```
-- MySQL
ALTER TABLE employees CHANGE old_column_name new_column_name VARCHAR(50);

-- PostgreSQL
ALTER TABLE employees RENAME COLUMN old_column_name TO new_column_name;

-- SQL Server
EXEC sp_rename 'employees.old_column_name', 'new_column_name', 'COLUMN';

-- Oracle
ALTER TABLE employees RENAME COLUMN old_column_name TO new_column_name;
```

**What are the different types of indexes in SQL?**

#### 1. Clustered Index

- Physical ordering of data matches index order
- One per table (usually on primary key)
- Faster for range queries

#### 2. Non-Clustered Index

- Separate structure pointing to data rows
- Multiple allowed per table
- Faster for specific lookups

#### 3. Unique Index

- Ensures uniqueness of indexed columns
- Automatically created for PRIMARY KEY and UNIQUE constraints

#### 4. Composite Index

- Index on multiple columns
- Order of columns matters for query optimization

#### 5. Partial Index

- Index on subset of rows (with WHERE condition)
- Saves space and improves performance

```
-- Creating different types of indexes
CREATE INDEX idx_employee_name ON employees(name);
CREATE UNIQUE INDEX idx_employee_email ON employees(email);
CREATE INDEX idx_emp_dept_salary ON employees(department, salary);
CREATE INDEX idx_high_salary ON employees(salary) WHERE salary > 50000;
```

## Data Types & Constraints

### What is the difference between CHAR and VARCHAR?

Aspect	CHAR	VARCHAR
<b>Storage</b>	Fixed-length	Variable-length
<b>Padding</b>	Right-padded with spaces	No padding
<b>Storage Efficiency</b>	Less efficient for short strings	More efficient
<b>Performance</b>	Slightly faster for fixed-size data	Slightly slower due to length calculation
<b>Use Case</b>	Status codes, country codes	Names, descriptions

```
-- CHAR example - always uses 10 bytes
name CHAR(10) -- 'John' stored as 'John      '

-- VARCHAR example - uses only needed bytes + length info
name VARCHAR(10) -- 'John' stored as 'John' (4 bytes + length)
```

### What is a PRIMARY KEY? What is an ALTERNATIVE KEY?

#### PRIMARY KEY:

- Uniquely identifies each record in a table
- Cannot contain NULL values
- Only one per table
- Automatically creates a unique index
- Cannot be changed once defined

#### ALTERNATIVE KEY (Candidate Key):

- Could serve as primary key but wasn't chosen
- Must be unique and not null
- Multiple alternative keys possible per table
- Often implemented using UNIQUE constraint

```
CREATE TABLE employees (
    emp_id INT PRIMARY KEY,           -- Primary Key
    ssn VARCHAR(11) UNIQUE NOT NULL, -- Alternative Key
    email VARCHAR(100) UNIQUE NOT NULL, -- Alternative Key
    name VARCHAR(50)
);
```

### How much storage does CHAR(10) take versus VARCHAR(10) ?

**CHAR(10):**

- Always uses exactly 10 bytes
- 'Hello' → 'Hello ' (10 bytes)

**VARCHAR(10):**

- Uses 1-2 bytes for length + actual data length
- 'Hello' → 1-2 bytes (length) + 5 bytes (data) = 6-7 bytes
- Maximum: 1-2 bytes + 10 bytes = 11-12 bytes

## What are the different constraints you can apply to a column?

1. **NOT NULL:** Prevents null values
2. **UNIQUE:** Ensures all values are unique
3. **PRIMARY KEY:** Combination of NOT NULL and UNIQUE
4. **FOREIGN KEY:** Links to another table's primary key
5. **CHECK:** Validates data against a condition
6. **DEFAULT:** Provides default value when none specified

```
CREATE TABLE employees (
    emp_id INT PRIMARY KEY,
    name VARCHAR(50) NOT NULL,
    email VARCHAR(100) UNIQUE NOT NULL,
    age INT CHECK (age ≥ 18 AND age ≤ 65),
    department VARCHAR(30) DEFAULT 'General',
    manager_id INT,
    FOREIGN KEY (manager_id) REFERENCES employees(emp_id)
);
```

## Normalization

### Explain the different normal forms in database normalization.

**1NF (First Normal Form):**

- Each column contains atomic (indivisible) values
- No repeating groups or arrays
- Each record is unique

**2NF (Second Normal Form):**

- Must be in 1NF
- All non-key attributes fully depend on the primary key
- Eliminates partial dependencies

**3NF (Third Normal Form):**

- Must be in 2NF
- No transitive dependencies

- Non-key attributes depend only on primary key

**BCNF (Boyce-Codd Normal Form):**

- Stricter version of 3NF
- Every determinant must be a candidate key

## Normalization Issues Example

**Problem Table:**

```
Orders(order_id, customer_name, customer_address, product1, price1, product2, price2)
```

**Issues:**

1. **1NF Violation:** Multiple products in single row
2. **Data Redundancy:** Customer info repeated
3. **Update Anomalies:** Changing customer address requires multiple updates
4. **Insert Anomalies:** Can't add customer without order

**Solution:**

```
-- Normalized structure
Customers(customer_id, customer_name, customer_address)
Products(product_id, product_name, price)
Orders(order_id, customer_id, order_date)
Order_Items(order_id, product_id, quantity)
```

## Query-Based Questions

### Basic Queries

**Find employees in HR department with salary above 50,000**

```
SELECT e.emp_id, e.name, e.salary
FROM employees e
JOIN departments d ON e.dept_id = d.dept_id
WHERE d.dept_name = 'HR'
AND e.salary > 50000;
```

### Aggregation & Grouping

**Find the number of employees in each department**

```
SELECT d.dept_name, COUNT(e.emp_id) as employee_count
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id
```

```
GROUP BY d.dept_id, d.dept_name
ORDER BY employee_count DESC;
```

## Find average salary excluding employees above company average

```
WITH company_avg AS (
    SELECT AVG(salary) as avg_salary
    FROM employees
)
SELECT d.dept_name, AVG(e.salary) as dept_avg_salary
FROM employees e
JOIN departments d ON e.dept_id = d.dept_id
CROSS JOIN company_avg ca
WHERE e.salary ≤ ca.avg_salary
GROUP BY d.dept_id, d.dept_name;
```

## Join Operations

### Find all employees and their managers

```
SELECT
    e.emp_id,
    e.name as employee_name,
    COALESCE(m.name, 'No Manager') as manager_name
FROM employees e
LEFT JOIN employees m ON e.manager_id = m.emp_id
ORDER BY e.emp_id;
```

## Highest Salary Queries

### Find the 3rd highest salary

```
-- Method 1: Using DENSE_RANK()
WITH salary_ranks AS (
    SELECT salary, DENSE_RANK() OVER (ORDER BY salary DESC) as rank
    FROM employees
)
SELECT DISTINCT salary
FROM salary_ranks
WHERE rank = 3;

-- Method 2: Using LIMIT/TOP with subquery
SELECT DISTINCT salary
FROM employees
```

```
ORDER BY salary DESC
LIMIT 1 OFFSET 2; -- PostgreSQL/MySQL

-- Method 3: Using correlated subquery
SELECT DISTINCT salary
FROM employees e1
WHERE 2 = (
    SELECT COUNT(DISTINCT salary)
    FROM employees e2
    WHERE e2.salary > e1.salary
);
```

## Find department with highest average salary

```
WITH dept_avg AS (
    SELECT
        d.dept_id,
        d.name as dept_name,
        AVG(e.salary) as avg_salary
    FROM departments d
    JOIN employees e ON d.dept_id = e.dept_id
    GROUP BY d.dept_id, d.name
)
SELECT dept_name, avg_salary
FROM dept_avg
WHERE avg_salary = (SELECT MAX(avg_salary) FROM dept_avg);
```

## Top N Records

### Find top 5 customers with highest total order amounts

```
SELECT
    customer_id,
    SUM(order_amount) as total_amount
FROM orders
GROUP BY customer_id
ORDER BY total_amount DESC
LIMIT 5;
```

### Find top 5 selling products in last 3 months

```
SELECT
    product_id,
    SUM(quantity) as total_sold
```

```
FROM sales
WHERE sale_date ≥ CURRENT_DATE - INTERVAL '3 months'
GROUP BY product_id
ORDER BY total_sold DESC
LIMIT 5;
```

## Advanced SQL Questions

### Window Functions

#### Rank employees by salary within each department

```
SELECT
    emp_id,
    name,
    dept_id,
    salary,
    RANK() OVER (PARTITION BY dept_id ORDER BY salary DESC) as salary_rank
FROM employees
ORDER BY dept_id, salary_rank;
```

#### Difference between RANK(), DENSE\_RANK(), and ROW\_NUMBER()

```
SELECT
    name,
    salary,
    ROW_NUMBER() OVER (ORDER BY salary DESC) as row_num,
    RANK() OVER (ORDER BY salary DESC) as rank_num,
    DENSE_RANK() OVER (ORDER BY salary DESC) as dense_rank_num
FROM employees
ORDER BY salary DESC;
```

/\*

Example Output:

name	salary	row_num	rank_num	dense_rank_num
Alice	100000	1	1	1
Bob	100000	2	1	1 -- Same salary
Charlie	90000	3	3	2 -- RANK skips 2, DENSE_RANK doesn't
David	90000	4	3	2
Eve	80000	5	5	3 -- RANK skips 4

\*/

#### Key Differences:

- **ROW\_NUMBER():** Always unique, sequential numbers



- **RANK()**: Same rank for ties, skips next rank(s)
- **DENSE\_RANK()**: Same rank for ties, no gaps in ranking

## Compare current month's sales with previous month using LAG/LEAD

```
SELECT
    month,
    sales_amount,
    LAG(sales_amount) OVER (ORDER BY month) as prev_month_sales,
    LEAD(sales_amount) OVER (ORDER BY month) as next_month_sales,
    sales_amount - LAG(sales_amount) OVER (ORDER BY month) as
month_over_month_change,
    ROUND(
        (sales_amount - LAG(sales_amount) OVER (ORDER BY month)) * 100.0 /
        LAG(sales_amount) OVER (ORDER BY month), 2
    ) as percent_change
FROM monthly_sales
ORDER BY month;
```

## Duplicate Handling

### Find duplicate records

```
-- Method 1: Using GROUP BY and HAVING
SELECT
    name, email, COUNT(*) as duplicate_count
FROM employees
GROUP BY name, email
HAVING COUNT(*) > 1;

-- Method 2: Using window functions
SELECT *
FROM (
    SELECT *,
        ROW_NUMBER() OVER (PARTITION BY name, email ORDER BY emp_id) as row_num
    FROM employees
) t
WHERE row_num > 1;
```

### Delete duplicate records keeping only one copy

```
-- Using CTE and ROW_NUMBER()
WITH duplicate_cte AS (
    SELECT *,
```

```

        ROW_NUMBER() OVER (
            PARTITION BY name, email
            ORDER BY emp_id -- Keep the one with smallest emp_id
        ) as row_num
    FROM employees
)
DELETE FROM employees
WHERE emp_id IN (
    SELECT emp_id
    FROM duplicate_cte
    WHERE row_num > 1
);

-- Alternative method using self-join
DELETE e1 FROM employees e1
INNER JOIN employees e2
WHERE e1.emp_id > e2.emp_id
    AND e1.name = e2.name
    AND e1.email = e2.email;

```

## Correlated Subqueries

Find employees who earn more than their department's average salary

```

SELECT emp_id, name, dept_id, salary
FROM employees e1
WHERE salary > (
    SELECT AVG(salary)
    FROM employees e2
    WHERE e2.dept_id = e1.dept_id
);

-- Alternative using window functions (more efficient)
WITH dept_avg AS (
    SELECT *,
        AVG(salary) OVER (PARTITION BY dept_id) as dept_avg_salary
    FROM employees
)
SELECT emp_id, name, dept_id, salary
FROM dept_avg
WHERE salary > dept_avg_salary;

```

## Query Performance Optimization

How would you optimize this query?

```
-- Original query
SELECT *
FROM orders o
JOIN customers c ON o.customer_id = c.customer_id
WHERE o.order_date ≥ '2024-01-01'
      AND o.status = 'completed'
      AND c.country = 'USA';
```

### Optimization strategies:

#### 1. Add indexes:

```
CREATE INDEX idx_orders_date_status ON orders(order_date, status);
CREATE INDEX idx_customers_country ON customers(country);
CREATE INDEX idx_orders_customer_id ON orders(customer_id);
```

#### 2. Select only needed columns:

```
SELECT o.order_id, o.order_date, o.order_amount, c.customer_name
FROM orders o
JOIN customers c ON o.customer_id = c.customer_id
WHERE o.order_date ≥ '2024-01-01'
      AND o.status = 'completed'
      AND c.country = 'USA';
```

#### 3. Filter early and use covering indexes:

```
-- Create covering index
CREATE INDEX idx_orders_covering ON orders(order_date, status, customer_id,
order_id, order_amount);

-- Rewrite query to filter orders first
SELECT o.order_id, o.order_date, o.order_amount, c.customer_name
FROM (
    SELECT customer_id, order_id, order_date, order_amount
    FROM orders
    WHERE order_date ≥ '2024-01-01' AND status = 'completed'
) o
JOIN customers c ON o.customer_id = c.customer_id
WHERE c.country = 'USA';
```

## What will you do if your query is running slow?

#### 1. Analyze Execution Plan:

```
EXPLAIN ANALYZE SELECT ... -- PostgreSQL
EXPLAIN EXECUTION PLAN SELECT ... -- Oracle
SET SHOWPLAN_ALL ON; SELECT ... -- SQL Server
```

## 2. Check for missing indexes:

- Look for table scans in execution plan
- Add indexes on WHERE, JOIN, and ORDER BY columns

## 3. Optimize WHERE clauses:

- Put most selective conditions first
- Avoid functions on indexed columns
- Use EXISTS instead of IN for large subqueries

## 4. Review JOIN operations:

- Ensure JOIN conditions use indexed columns
- Consider JOIN order for multiple tables

## 5. Update table statistics:

```
ANALYZE TABLE table_name; -- MySQL
UPDATE STATISTICS table_name; -- SQL Server
```

## 6. Consider query rewriting:

- Replace correlated subqueries with JOINS
- Use window functions instead of multiple GROUP BYs
- Break complex queries into simpler parts

# Conditional Aggregation

## Count employees in each department who joined in 2024

```
-- With GROUP BY
SELECT
    d.dept_name,
    COUNT(CASE WHEN YEAR(e.hire_date) = 2024 THEN 1 END) as new_hires_2024,
    COUNT(e.emp_id) as total_employees
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id
GROUP BY d.dept_id, d.dept_name;

-- Without GROUP BY (using window functions)
SELECT DISTINCT
    d.dept_name,
    COUNT(CASE WHEN YEAR(e.hire_date) = 2024 THEN 1 END)
        OVER (PARTITION BY d.dept_id) as new_hires_2024,
```

```
COUNT(e.emp_id) OVER (PARTITION BY d.dept_id) as total_employees
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id;
```

## Common Table Expressions (CTE)

### Rewrite subquery using CTE

```
-- Original with subquery
SELECT dept_name,
       (SELECT COUNT(*)
        FROM employees e
        WHERE e.dept_id = d.dept_id) as emp_count
FROM departments d;

-- Rewritten with CTE
WITH dept_employee_count AS (
    SELECT
        dept_id,
        COUNT(*) as emp_count
    FROM employees
    GROUP BY dept_id
)
SELECT
    d.dept_name,
    COALESCE(dec.emp_count, 0) as emp_count
FROM departments d
LEFT JOIN dept_employee_count dec ON d.dept_id = dec.dept_id;
```

## Date-Based Queries

### Find all orders from the last 3 months with rouge date handling

```
SELECT *
FROM orders
WHERE order_date ≥ CURRENT_DATE - INTERVAL '3 months'
      AND order_date ≤ CURRENT_DATE -- Handle future dates
      AND order_date ≥ '1900-01-01' -- Handle unrealistic past dates
      AND order_date IS NOT NULL    -- Handle NULL dates
ORDER BY order_date DESC;

-- More robust version with CASE statement
SELECT *,
       CASE
           WHEN order_date IS NULL THEN 'Missing Date'
```

```

        WHEN order_date > CURRENT_DATE THEN 'Future Date'
        WHEN order_date < '1900-01-01' THEN 'Invalid Past Date'
        ELSE 'Valid Date'
    END as date_status
FROM orders
WHERE order_date ≥ CURRENT_DATE - INTERVAL '3 months'
    AND order_date ≤ CURRENT_DATE
    AND order_date ≥ '1900-01-01'
    AND order_date IS NOT NULL;

```

## Transaction Management

### What are ACID properties?

**ACID** ensures database transactions are processed reliably:

1. **Atomicity:** Transaction is all-or-nothing
  - Either all operations succeed or all fail
  - No partial transactions
2. **Consistency:** Database remains in valid state
  - All constraints and rules are maintained
  - Data integrity is preserved
3. **Isolation:** Concurrent transactions don't interfere
  - Transactions appear to run sequentially
  - Prevents dirty reads, phantom reads
4. **Durability:** Committed changes are permanent
  - Survive system crashes
  - Changes are written to persistent storage

```

-- Example transaction demonstrating ACID
BEGIN TRANSACTION;

    UPDATE accounts SET balance = balance - 100 WHERE account_id = 1;
    UPDATE accounts SET balance = balance + 100 WHERE account_id = 2;

    -- Check if both accounts have sufficient balance
    IF (SELECT balance FROM accounts WHERE account_id = 1) < 0
        ROLLBACK; -- Atomicity: undo all changes
    ELSE
        COMMIT;   -- Durability: make changes permanent
END TRANSACTION;

```

## NULL Handling

### Handle null values in salary calculation

```
-- Different approaches to handle NULLs
SELECT
    emp_id,
    name,
    salary,

    -- Replace NULL with 0
    COALESCE(salary, 0) as salary_with_zero,

    -- Replace NULL with average salary
    COALESCE(salary, (SELECT AVG(salary) FROM employees WHERE salary IS NOT NULL))
as salary_with_avg,

    -- Use CASE statement
    CASE
        WHEN salary IS NULL THEN 0
        ELSE salary
    END as salary_case,

    -- Calculate bonus (10% of salary, 0 if NULL)
    COALESCE(salary * 0.10, 0) as bonus,

    -- Total compensation
    COALESCE(salary, 0) + COALESCE(bonus_amount, 0) as total_compensation
FROM employees;

-- Aggregate functions with NULL handling
SELECT
    dept_id,
    COUNT(*) as total_employees,
    COUNT(salary) as employees_with_salary,    -- Excludes NULLs
    AVG(salary) as avg_salary,                  -- Ignores NULLs
    AVG(COALESCE(salary, 0)) as avg_salary_with_zeros
FROM employees
GROUP BY dept_id;
```

## ⋮ Data Cleaning Queries

### Handle duplicate values and standardize data

```
-- Clean and standardize customer data
WITH cleaned_customers AS (
    SELECT
        customer_id,
        -- Standardize names
```

```

INITCAP(TRIM(LOWER(customer_name))) as clean_name,

-- Clean phone numbers
REGEXP_REPLACE(phone, '[^0-9]', '') as clean_phone,

-- Standardize email
LOWER(TRIM(email)) as clean_email,

-- Clean addresses
INITCAP(TRIM(REGEXP_REPLACE(address, '\s+', ' '))) as clean_address,

-- Handle NULL values
COALESCE(city, 'Unknown') as city,
UPPER(COALESCE(state, 'XX')) as state
FROM customers
WHERE email IS NOT NULL
      AND email LIKE '%@%' -- Basic email validation
)
SELECT * FROM cleaned_customers;

-- Remove duplicates based on business rules
WITH ranked_customers AS (
    SELECT *,
           ROW_NUMBER() OVER (
               PARTITION BY clean_email
               ORDER BY
                   CASE WHEN phone IS NOT NULL THEN 1 ELSE 2 END,
                   customer_id
           ) as rn
    FROM cleaned_customers
)
SELECT * FROM ranked_customers WHERE rn = 1;

```

## Query Performance Best Practices

### Indexing Strategy

- 1. Primary Key Index:** Automatically created
- 2. Foreign Key Indexes:** Always index foreign key columns
- 3. Covering Indexes:** Include all columns needed by query
- 4. Composite Indexes:** Order columns by selectivity (most selective first)

```

-- Good composite index order
CREATE INDEX idx_orders_date_status_customer
ON orders(order_date, status, customer_id);

```



```
-- Include frequently selected columns
CREATE INDEX idx_orders_covering
ON orders(order_date, status)
INCLUDE (customer_id, order_amount);
```

## Query Writing Best Practices

1. Select only needed columns
2. Use appropriate JOINS
3. Filter early in subqueries
4. Avoid functions on indexed columns in WHERE clause
5. Use EXISTS instead of IN for large result sets
6. Consider using UNION ALL instead of UNION when duplicates are acceptable

```
-- Bad: Function on indexed column
SELECT * FROM orders WHERE YEAR(order_date) = 2024;

-- Good: Range condition
SELECT * FROM orders
WHERE order_date ≥ '2024-01-01'
AND order_date < '2025-01-01';
```